

Sensor Complementarity and Operating-Point Proxies in Acoustic Emission and Passive Ultrasound Regression of the ISO 281 Viscosity Ratio

Frederik Appel Vardinghus-Nielsen

Department of Electrical and Computer Engineering, Aarhus University, Denmark
frederikavn@gmail.com

Draft submitted for peer review. Do not cite or distribute without permission.

Keywords: acoustic emission; lubrication condition monitoring; viscosity ratio κ ; ball bearing; signal processing; sensor fusion

Abstract

Inadequate lubrication causes an estimated 30–80% of rolling-element bearing failures, yet online, non-invasive lubrication condition monitoring (LCM) remains unsolved: lubrication-relevant signal content is weak and strongly confounded with operating conditions. This paper presents a controlled dual-channel study in which wideband acoustic emission (AE) and heterodyned passive ultrasound (US) are acquired simultaneously from a ball bearing across staircase speed sweeps (140–2977 rpm) and 13 temperature blocks (33–100 °C), spanning the boundary, mixed, and full-film regimes, with the ISO 281 viscosity ratio κ as a continuous, physics-grounded regression target. Because κ is determined by the operating point, the central methodological question is what a feature- κ regression actually senses: a two-stage decomposition shows that acoustically inferred operating points alone recover $R^2 = 0.826$ of the direct model’s 0.834, while within-speed-step analysis isolates a genuine temperature-mediated lubrication channel (20–500 kHz mobility, $\rho = -0.56$ at fixed drive state). Evaluation uses visit-grouped data splits that remove the near-duplicate leakage of windowed acquisition; recorded content above the AE coupler’s 1 MHz low-pass is excluded from analysis. Under this protocol, gradient boosting attains hold-out $R^2 = 0.834$ (AE) and 0.653 (US); fusing the channels raises performance to 0.908 ($\Delta\text{RMSE} = 0.0281$, significant on all tests), answering the long-argued but previously untested question of AE–US complementarity. All results derive from repeated grouped cross-validation with corrected significance tests.

Contents

1	Introduction	5
2	Background	9
2.1	Lubrication regimes and the viscosity ratio κ	9
2.2	Physical origins of AE and passive ultrasound signals	9
3	Experimental Setup	10
3.1	Test rig and bearing	10
3.2	Operating conditions	11
3.3	Sensor configuration and data acquisition	12
4	Signal Processing and Feature Extraction	14
4.1	Waveforms and spectrograms	14
4.2	Sub-band decomposition	15
4.3	Sub-band signal validation	16
4.4	Feature extraction	16
4.5	Candidate feature sets	17
5	Feature Analysis	18
5.1	Two-stage feature selection	18
5.2	Correlation with κ	19
5.3	Redundancy structure	20
5.4	Operating-condition dependence	21
5.5	Low-dimensional structure and regime separability	22
6	Regression Modelling and Statistical Evaluation	23
6.1	Experimental design	23
6.1.1	Data sets and splits	23
6.1.2	Model families	23
6.2	Repeated nested cross-validation	23
6.3	Statistical significance testing	25
6.4	Results	26
7	Discussion	28
7.1	AE vs. US	28
7.2	Fusion gain	29
7.3	Interpreting the amplitude-confounding result	29
7.4	Implications for the κ regression framework	30
7.5	Limitations	30

8	Conclusion	31
A	Full Feature Correlation Rankings	32

Nomenclature

Symbols

κ	ISO 281 viscosity ratio, ν/ν_1
ν	kinematic viscosity at operating temperature [mm ² s ⁻¹]
ν_1	reference viscosity for full film formation [mm ² s ⁻¹]
n	rotational speed [rpm]
d_{pw}	mean pitch diameter [mm]
N	samples per sweep
x_n	voltage sample n
S_k, f_k	one-sided FFT magnitude and frequency at bin k
f_c	spectral centroid [Hz]
σ_w	magnitude-weighted spectral bandwidth [Hz]
ρ, r	Spearman and Pearson correlation coefficients
R, k	cross-validation repetitions and folds

Abbreviations

AE	acoustic emission
CV	cross-validation
EHL	elastohydrodynamic lubrication
EMI	electromagnetic interference
HO	hold-out (test set)
LCM	lubrication condition monitoring
MAE	mean absolute error
RMSE	root mean square error
SHAP	SHapley Additive exPlanations
US	passive ultrasound
VFD	variable-frequency drive
VIF	variance inflation factor

1 Introduction

Rolling element bearings are among the most prevalent components in industrial machinery, and their reliable operation depends critically on adequate lubrication. Lubrication failure—through starvation, viscosity degradation, or contamination—increases metal-to-metal asperity contact, accelerates surface fatigue and wear, raises operating temperature, and ultimately causes premature failure. Improper lubrication is estimated to account for between 30% and 80% of all rolling element bearing failures [1–3].

Despite this, lubrication condition monitoring (LCM) has received far less attention than vibration-based fault detection and remaining-useful-life estimation, which are now mature and widely deployed [4, 5]. The lubrication signal is harder to exploit: its acoustic and vibratory signatures are weaker than the impulsive signatures of spalling or raceway cracks and are strongly entangled with speed and load, so operating-condition effects readily mask the lubrication-dependent component [5–7]. Within bearing fault diagnosis, this entanglement has motivated a substantial recent literature on diagnosis under varying and unseen working conditions [8–10]; for LCM the problem is sharper still, because the lubrication target itself is derived from the operating point—a confound this paper addresses with an explicit diagnostic decomposition (Section 5.4). Yet accurate online LCM would enable condition-based relubrication—applying lubricant when and in the quantity actually needed, rather than on fixed schedules that systematically under- and over-lubricate [11].

The dominant bearing-monitoring approach, vibration analysis, detects macroscopic faults such as raceway spalling and rolling-element damage through characteristic-frequency analysis. Classical characteristic-frequency analysis is, however, poorly matched to the subtle, non-stationary tribological changes of film formation and breakdown, which lack the stable periodic signatures of structural faults [5, 12]; recent data-driven work nonetheless recovers lubrication-state information from vibration, typically from high-frequency broadband energy rather than from characteristic frequencies [13]. Temperature monitoring and oil analysis add complementary information but are indirect, subject to sampling delay, or require access to the lubricant reservoir, and none provides a direct, continuous, non-invasive indicator of the instantaneous lubrication regime under operating conditions [11, 14].

Passive, non-invasive sensing, e.g. acoustic emission (AE) and passive ultrasound (US), is practically attractive for industrial LCM because it requires no bearing modification, no direct access to the lubricant, and can be retrofitted to bearings in service. AE captures the broadband stress waves emitted by asperity contact and microcrack events, while passive US underlies the threshold-based relubrication tools already in routine industrial use, in which the broadband ultrasonic level is compared against a baseline and relubrication is triggered once it rises by a set margin [1, 15, 16]. The physical origins of

both signals are detailed in Section 2. Despite the practical maturity of threshold-based US and the demonstrated usefulness of AE, the two channels’ signal-processing behaviour has not been characterised simultaneously within a single regression framework against a physics-grounded lubrication reference.

The sensitivity of AE to lubrication regime is well established. Miettinen et al. [17] showed that AE pulse count and RMS distinguish boundary, mixed, and full-film regimes and respond to base-oil viscosity and thickener content in grease-lubricated bearings. Yoshioka and Shimizu [18] combined AE and vibration monitoring on a grease-lubricated deep-groove ball bearing using envelope demodulation and statistical AE parameters. Tandon et al. [19] found AE and shock-pulse sensors to outperform accelerometers for detecting grease contamination in a ball bearing—the scenario closest to the present study. More recently, Cornel et al. [12] showed that AE detects lubrication-sensitive surface damage such as micro-pitting and smearing up to eight hours earlier than vibration monitoring, Wang et al. [20] developed AE fingerprint features with dynamic thresholds for bearing fault detection, and Hou and Zhang [21] derived amplitude- and energy-centre features that classify lubrication status across speeds and viscosity levels, though without a continuous regression target, and with no variation of temperature.

On the signal-processing side, Renhart et al. [22] applied octave filter-bank decomposition with RMS–friction features to characterise AE tribological events, and Marticorena and García Peyrano [6] showed that spectral correlation of AE resolves lubrication-regime changes under varying load, demonstrating that kinematic and acoustic features can partially decouple operating-condition confounds. The physical link between film adequacy and the AE signature has been formalised by Liu et al. [23] and Zhang et al. [7], whose coupled vibration–AE models tie emission intensity and frequency content to EHL film thickness and the rate of asperity contact. For the passive-ultrasound channel, Jakobsen et al. [1, 16] showed that broadband acoustic features carry measurable information about the ISO 281 viscosity ratio κ .

A recurring limitation of data-driven LCM is its reliance on discrete fault labels or qualitative lubrication states (e.g. [13]) rather than a physically interpretable, continuous target. The ISO 281 viscosity ratio κ removes this limitation: it is the ratio of the lubricant’s actual kinematic viscosity at operating temperature to the minimum viscosity required for full EHL film formation at the current speed and bearing geometry, mapping the instantaneous operating point to a standardised lubrication-adequacy index (Section 2). A model trained to predict κ implicitly learns the joint dependence on speed, temperature, and lubricant properties, and its output is directly interpretable against established regime boundaries. Despite this, κ has seldom been used as a regression objective in acoustic condition monitoring.

The clearest precedent is the work of Jakobsen et al. [1, 16], who established κ regression for passive bearing acoustics, predicting κ from low-cost MEMS microphone and

accelerometer features using LASSO and neural networks. The present study adopts their κ target and extends it in three respects: it adds a wideband AE channel, analysed up to the 1 MHz limit of the AE coupler’s band-pass, acquired simultaneously with the passive US channel, extracts frequency sub-band features from both signals, and tests whether the two modalities are complementary. This last point is an open question: AE–US complementarity is argued on physical grounds in reviews [5, 24] but has not been quantified experimentally, and it carries direct practical weight, since a second sensor adds cost and complexity and is justified only if it supplies non-redundant information. No prior study has simultaneously acquired AE and passive US from a single ball bearing and quantitatively assessed their joint value for κ regression.

This study deliberately fixes the bearing, lubricant, and load, varying only rotational speed and temperature. Holding the mechanical configuration constant isolates the lubrication-dependent acoustic response from the bearing-to-bearing variability and degradation processes that would otherwise confound it—a prerequisite for establishing whether AE and passive US carry κ -relevant information at all. The resulting limits on generalisation across bearing types, lubricants, and load regimes are genuine, but not addressed here.

Contributions

This paper makes the following contributions:

1. **A confound-decomposition methodology for acoustic LCM.** Because κ is determined by the operating point, any feature– κ regression risks acting as an operating-point soft sensor. A two-stage decomposition (features $\rightarrow$ RPM/temperature $\rightarrow$ ISO 281 mapping, recovering $R^2 = 0.826$ of the direct model’s 0.834) quantifies the proxy share, while within-RPM-step correlation analysis of 20–500 kHz mobility (median $\rho = -0.56$ with temperature at fixed drive state) isolates the genuine lubrication-channel sensitivity.
2. **The AE–US complementarity question is answered quantitatively.** This is the first study to simultaneously acquire wideband AE and heterodyned passive US from a single ball bearing and assess their joint value: fusing the channels reduces hold-out RMSE from 0.109 to 0.081 ($\Delta\text{RMSE} = 0.0281$), significant under repeated grouped cross-validation and all hold-out tests, with the extended 20–100 kHz US band contributing selected features to the combined model. AE outperforms US on the hold-out set, while fold-level inference between the single channels is inconclusive under group-respecting cross-validation.
3. **A leakage-aware, instrumentation-honest evaluation protocol.** Recorded content above the 1 MHz low-pass of the AE coupler (Kistler Type 5125C)—present

in the digitised signal despite that hardware filter, and of undetermined origin—is excluded from analysis. Visit-grouped data splits remove the near-duplicate leakage inherent to windowed acquisition of slowly varying protocols. Under matched random-split protocols, performance equals prior κ -regression work; the stricter visit-grouped figures are, to the authors’ knowledge, the first reported for this task.

4. **A dual-channel dataset and reproducible evaluation protocol.** 8 504 sweeps spanning the boundary through full-film regimes ($\kappa \in [0.04, 1.64]$) with an analytically derived ISO 281 κ reference, evaluated under repeated nested cross-validation (25 outer scores) with Nadeau–Bengio-corrected and Holm–Bonferroni-adjusted inference.

The remainder of the paper is organised as follows. Section 2 reviews the theory of κ and the physical origins of AE and US signals. Section 3 describes the experimental setup and data collection. Section 4 presents the signal-processing and feature-extraction pipeline, including the characterisation of the frequency sub-bands. Section 5 reports the feature analysis and the operating-condition confound decomposition. Section 6 presents the regression modelling and statistical evaluation. Section 7 discusses the findings, and Section 8 concludes.

The overall pipeline of the experiment, signal processing, modelling, and evaluation is shown in Figure 1.




Figure 1: Overview of the experimental and modelling pipeline. AE and passive US signals are acquired simultaneously from the test rig and decomposed into broadband and sub-band features. Three sensor configurations (AE-only, AE+US fusion, US-only) are evaluated across three regression model families under repeated nested cross-validation, with the viscosity ratio κ as the regression target.

Align with final data availability decision (full deposit, features+code only, or on request).

2 Background

2.1 Lubrication regimes and the viscosity ratio κ

Rolling element bearing contacts operate in one of three lubrication regimes depending on whether the lubricant film is sufficient to fully separate the rolling surfaces. In the full-film elastohydrodynamic (EHL) regime the surfaces are completely separated and asperity contact is suppressed; in the mixed regime the film is partially formed and intermittent asperity contact occurs; in the boundary regime the film is inadequate and metal-to-metal contact dominates. The ISO 281 standard characterises these regimes through the viscosity ratio κ , with boundaries at $\kappa < 0.5$ (boundary), $0.5 \leq \kappa < 1$ (mixed), and $\kappa \geq 1$ (full film).

The ISO 281 standard defines the viscosity ratio

$$\kappa = \frac{\nu}{\nu_1(n, d_{pw})}, \quad (1)$$

where ν [$\text{mm}^2 \text{s}^{-1}$] is the actual kinematic viscosity of the lubricant at operating temperature, and ν_1 is the minimum kinematic viscosity required for full EHL at the bearing's rotational speed n [rpm] and mean pitch diameter d_{pw} [mm] [25]. The reference viscosity is given by

$$\nu_1 = \begin{cases} 45\,000 \, n^{-0.83} d_{pw}^{-0.5} & n < 1000 \text{ rpm}, \\ 4\,500 \, n^{-0.5} d_{pw}^{-0.5} & n \geq 1000 \text{ rpm}. \end{cases} \quad (2)$$

The operating viscosity ν is computed from the ASTM D341 Walther viscosity–temperature equation fitted to the lubricant's 40 °C and 100 °C reference values, with a standard value of $\lambda = 0.7$. Since ν_1 decreases with increasing rotational speed—faster rotation enhances hydrodynamic film entrainment, requiring less viscosity to achieve full EHL— κ captures the combined effect of speed, temperature, and lubricant grade on lubrication adequacy at a given operating point. Importantly, κ assumes nominally intact lubricant: the Walther equation reflects the properties of fresh grease, so κ does not encode viscosity changes arising from lubricant degradation, contamination, or starvation. Extending the regression framework to track such changes would require either continuous viscosity updates or an alternative target variable.

2.2 Physical origins of AE and passive ultrasound signals

Acoustic emission (AE) refers to transient elastic stress waves generated by sudden energy releases in a material: asperity contact and micro-fracture are the dominant sources in rolling bearings [17, 22]. Under boundary and mixed lubrication, intermittent metal-to-metal asperity contacts generate broadband stress waves that propagate through the bearing structure and are detected by piezoelectric AE sensors coupled to the outer hous-

ing [12, 22]. The emission rate, amplitude, and frequency content of these events are all sensitive to the contact severity and hence to the lubrication regime [12, 17]. Full-film EHL suppresses asperity contact, reducing AE activity; the transition through the mixed regime is associated with intermittent, structurally distinct emission events. AE sensors for bearings respond from tens of kHz into the MHz range and require correspondingly high sampling rates to capture the full frequency content.

Passive ultrasound sensing detects structure-borne acoustic energy above the audible range. As a modality it is not intrinsically narrowband—ultrasonic emissions span from roughly 20 kHz into the MHz range, overlapping much of the AE band—but the instruments most widely deployed for bearing monitoring operate in a narrow band around 38–40 kHz. In this band the sensor amplitude is modulated by the contact-stiffness and viscous-damping changes associated with lubricant film formation: as the film degrades and asperity contact increases, the ultrasonic level at the bearing housing rises [15]. Commercial instruments (UE Systems, Sonotec) exploit this by comparing the level against a reference baseline and triggering relubrication when a defined threshold is exceeded, making threshold-based passive US the most widely deployed non-invasive LCM method in routine industrial use. The passive US probe used in this study follows the same design: it senses around 38–40 kHz and heterodynes (down-mixes) that band to an audio-frequency baseband (Section 3). In practice, however, the recorded signal also contains energy in bands above the nominal baseband. As this higher-frequency content may nonetheless carry lubrication-relevant information, it is retained in the analysis rather than discarded, without attributing it to a specific physical origin. However, the level is also sensitive to rotational speed and load independently of lubrication condition, which limits the reliability of threshold-based approaches under varying operating conditions.

Check
that this
is correct

The two modalities thus respond to the same underlying phenomenon—increasing asperity contact as the lubricant film degrades—but through different physical mechanisms and at different frequency scales: AE captures individual high-frequency transient events (≥ 50 kHz) from discrete micro-fractures and asperity collisions, while passive US, after the probe heterodynes its 38–40 kHz sensing band to an audio-frequency baseband, reflects the lower-frequency bulk acoustic envelope comprising structural resonances, friction noise, and the aggregate effect of contact activity, together with residual energy above the baseband whose information content is assessed empirically. Both channels therefore carry information about κ , and the question of whether they provide complementary or redundant lubrication-relevant information is addressed experimentally in Sections 3–6.

3 Experimental Setup

3.1 Test rig and bearing

Include
picture
of test
rig

The test rig is a custom-built bearing test stand at Aarhus University (Herning, Denmark). A single deep-groove ball bearing (model SYJ25, bore 25 mm, pitch diameter $d_{pw} = 38.0$ mm, 9 chrome steel rolling elements, dynamic load rating 14000 N) is driven by a variable-frequency electric motor controlled through a variable-frequency drive (VFD). The bearing is loaded radially by a static weight and lubricated by a defined quantity of Keratech 22 oil (kinematic viscosity $\nu_{40} = 22$ cSt, $\nu_{100} = 4.1$ cSt) .

Note on VFD noise: Harmonic artefacts from the VFD switching frequency are present at approximately 5 kHz (nominal; the measured comb spacing is ≈ 4.7 kHz) and its multiples in the low-frequency range of both sensor signals. These are partially filtered out

3.2 Operating conditions

The experiment follows a protocol which varies the operating conditions through changes in rotational speed (rpm) and housing temperature ($^{\circ}\text{C}$). This consists of a staircase sweep from 100 to 3000 RPM and back down to 100 RPM in steps of 100 RPM with a 60-second hold at each step. This staircase is repeated for each of 13 temperature blocks (40–100 $^{\circ}\text{C}$ in 5 $^{\circ}\text{C}$ increments) where temperature is controlled with . The block temperature set-point is changed between staircases rather than during the high-speed hold: after each downsweep the rig is held at the low-rpm bottom of the staircase (≈ 100 rpm) while the housing is brought to the next set-point and allowed to stabilise thermally. The recorded telemetry contains no true standstill (RPM = 0) sweeps; the only near-stationary data is a single warm-up segment (37–43 $^{\circ}\text{C}$) acquired before the first staircase, which is excluded from analysis (Section 3.3). The hold at the 3000 rpm staircase peak is substantially longer than the ordinary 60-second steps (approximately 20 minutes), so this operating point is strongly over-represented in the data: it accounts for roughly a third of all sweeps. The applied radial load is a static dead-weight, held constant throughout the experiment at approximately 61 kg (≈ 600 N). The total duration of the experiment is approximately 14 hours. The measured operating conditions span 140–2977 rpm and 33–100 $^{\circ}\text{C}$; the measured temperature range extends below the commanded 40 $^{\circ}\text{C}$ floor (down to ≈ 33 $^{\circ}\text{C}$), which we tentatively attribute to the inherent undershoot and instability of the temperature control rather than to any commanded set-point.

The lubrication conditions in the bearing are not expected to change during the experiment in other ways than what is entailed by rpm and temperature and the effect of these upon the lubricant. It is not expected or assumed that any degradation, starvation or contamination of the lubricant happens.

The operating conditions are controlled through . The measured operating conditions space, distribution and time variation are seen in Figures 2 and 3. The κ value for each sweep is computed according to Equations (1) and (2) using the per-sweep measured RPM

What is the amount?

Do we need NLGI consistency grade?

How is temperature controlled and measured?

Explain the process and expected outcome.

How is RPM controlled

With what?

The original protocol described a 2-min RPM = 0 standstill at each block start,

and temperature, with ν evaluated via the ASTM D341 Walther viscosity–temperature equation.




Figure 2: Left: space of the measured operating conditions (RPM and temperature) during the experiment. Each point corresponds to a 20ms window of sampling at a given RPM and temperature. The colour of each point corresponds to the derived κ value for that sweep, with warmer colours indicating higher κ values. Right: distribution of the derived κ values.




Figure 3: Measured operating conditions (RPM and temperature) for each sweep during the experiment. Each point corresponds to a 20ms window of sampling at a given RPM and temperature. The colour of each point corresponds to the derived κ value for that sweep, with warmer colours indicating higher κ values.

3.3 Sensor configuration and data acquisition

Two sensors simultaneously acquire signals from the bearing housing through a shared oscilloscope which samples at a shared sample rate of 12.5 MHz:

Consider
invert-
ing the
colors.
Green
should
be bet-
ter,
right?

Acoustic emission sensor: A Kistler Piezotron Type 8152C1 wideband piezoelectric AE sensor with integral impedance converter, bonded to the bearing housing. The sensor is ground-isolated, rated to 165°C, and has a nominal ± 10 dB frequency range of 100–900 kHz (sensitivity 48 dB_{ref} 1 V/(m/s)). Its output is conditioned by a Kistler Type 5125C AE Piezotron coupler, which powers the sensor and applies a Butterworth band-pass filter (50 kHz high-pass, 1 MHz low-pass) to the raw AE signal before it is digitised. This *coupler* band-pass, not the sensor element, sets the 1 MHz upper analysis limit used throughout the paper.

Ultrasound sensor:

Temperature sensor:

RPM sensor:

Load sensor:

The shared sample rate and DAQ setup ensures that AE and US signals share timestamps.

Acquisition windows (sweeps). Data are acquired as short windows rather than continuously: every ~ 5.9 s the oscilloscope captures a 20 ms window ($N \approx 2.5 \times 10^5$ samples per channel) containing time-synchronous AE and US waveforms, together with the mean RPM and housing temperature over the window. Throughout this paper, one such window is referred to as a *sweep*, and a sweep is the unit of observation for all feature extraction, model training, and evaluation. After excluding an initial stationary warm-up segment recorded before the first staircase, the protocol yields 8 739 raw sweeps, of which 235 with measured RPM above the 3000 rpm protocol maximum are discarded, leaving 8 504 sweeps per channel (median 10 sweeps per staircase hold). The resulting κ distribution spans [0.04, 1.64] with mean 0.51 (Figure 2).

AE sensor/coupler
bullets
written
from the
Kistler
datasheets
in
docs/.
Re-
maining
bullets
(US
probe,
temper-
ature,
RPM,
load)
await
collab-
orator
input.
Temper-
ature
via an
Omron
con-
troller?

4 Signal Processing and Feature Extraction

4.1 Waveforms and spectrograms




Figure 4: Waveforms of AE and UL signals from the same sweeps for different values of κ . AE shows clear pulses of energy from asperity contacts which modulate the signal afterwards. UL shows a more regular, lower frequency, non-stationary waveform.

Spectrograms of the raw signals, ordered chronologically by sweep, are shown in Figure 5 to illustrate the frequency content and its variation with operating conditions across the protocol.




Figure 5: Spectrograms of the raw AE and US ordered by sweep number (chronological). The dashed lines show the sub-bands which are later used in feature generation.

4.2 Sub-band decomposition

The AE and US signals are decomposed into frequency sub-bands prior to feature extraction, since the information on lubrication conditions contained in the signals is expected to be localised in sub-bands rather than wholly contained in the broadband signals [7, 22]. Adaptive decompositions such as symplectic geometry mode decomposition offer data-driven alternatives for isolating informative components [26]; fixed, physically motivated bands are preferred here so that features remain directly comparable across sweeps and interpretable against the sensing physics. For the AE channel, two sub-bands are extracted in addition to the broadband signal: 20–500 kHz and 500 kHz–1 MHz which contain bands of spectral energy which upon visual inspection of Figure 5 vary with the operating conditions. Content above the 1 MHz coupler low-pass (Section 3.3) is excluded from the feature set. For the US channel, the signal spans the probe’s 0–100 kHz acquisition band, and three sub-bands are extracted in addition to the broadband signal: 0–10 kHz and 10–20 kHz within the nominal heterodyned baseband (Section 2.2), and 20–100 kHz capturing the above-baseband content whose information value is assessed empirically. The sub-bands are summarised in Table 1. All decompositions use 5th-order Butterworth filters applied forward–backward (`sosfiltfilt`), giving zero phase distortion and a 10th-order effective magnitude response.

Table 1: Frequency sub-bands used for signal pre-filtering before feature extraction.

Channel	Sub-band	Filter type	Frequency range
AE	Broadband	–	0–1 MHz
AE	Sub-band 1	Band-pass	20–500 kHz
AE	Sub-band 2	Band-pass	500 kHz–1 MHz
US	Broadband	Low-pass	0–100 kHz
US	Sub-band 1	Band-pass	0–10 kHz
US	Sub-band 2	Band-pass	10–20 kHz
US	Sub-band 3	Band-pass	20–100 kHz

4.3 Sub-band signal validation

The retained sub-bands carry bearing-generated content rather than drive EMI. The VFD injects a switching comb at ≈ 4.7 kHz and its harmonics (Section 3.1); these lie below the 20 kHz lower edge of the retained AE bands and so cannot masquerade as bearing signal within them. More directly, total band power rises with rotational speed in both bands— $\rho = +0.91$ for 20–500 kHz (at 45–55 °C) and $\rho = +0.69$ for 500 kHz–1 MHz (Figure 6)—as expected for mechanically generated contact and rolling noise [7, 17], rather than the fixed-frequency behaviour of an electromagnetic-interference source.




Figure 6: Total power of the two retained AE sub-bands as a function of RPM at fixed temperature windows. Both bands rise with speed, consistent with mechanically generated emission.

4.4 Feature extraction

Features are computed from the voltage signal of each sweep–sensor pair, for the broadband signal and for each sub-band. In total 14 features are computed, comprising 8 time-domain and 6 frequency-domain statistics, as listed in Tables 2 and 3. The set is deliberately canonical: commonly used companions such as standard deviation, variance, peak, impulse factor, and RMS frequency are exact transforms or products of the retained features (e.g. $\text{std} = \text{RMS}$ for zero-mean band-filtered signals; $\text{peak} = \text{crest} \times \text{RMS}$; $\text{impulse} = \text{crest} \times \text{shape}$; $f_{\text{RMS}}^2 = f_c^2 + \sigma_w^2$) and are therefore omitted, so that no candi-

date feature is a deterministic function of another. The frequency-centroid information that RMS frequency would carry is instead represented by the retained Hjorth mobility, its time-domain analogue (mobility $\approx 2\pi f_{\text{RMS}}$): computed from $\text{Var}(x')/\text{Var}(x)$ rather than from the spectral moments, it is not an exact function of f_c and σ_w . All retained features (without the specific sub-bands) have previously been applied to signals for bearing condition monitoring and LCM [1, 27]; the Hjorth parameters originate from EEG analysis [28].

Name	Formula	Description
RMS	$\sqrt{N^{-1} \sum_n x_n^2}$	Overall signal energy in the band; grows with the rate and intensity of contact activity.
Skewness	$N^{-1} \sum_n [(x_n - \bar{x})/\sigma]^3$	Asymmetry of the amplitude distribution; departs from zero when excursions are predominantly one-sided.
Kurtosis	$N^{-1} \sum_n [(x_n - \bar{x})/\sigma]^4$	Impulsiveness; large for sparse, sharp transients and near 3 for Gaussian background noise.
Crest factor	$x_{\text{peak}}/x_{\text{RMS}}$	Prominence of isolated peaks above the background level; elevated by sharp, intermittent events.
Shape factor	$x_{\text{RMS}} / (N^{-1} \sum_n x_n)$	Waveform peakiness independent of amplitude; larger for spiky signals than for smooth, sinusoidal ones.
Margin factor	$x_{\text{peak}} / (N^{-1} \sum_n \sqrt{ x_n })^2$	Sensitivity to extreme excursions, weighting large peaks more strongly than the crest factor.
Hjorth mobility	$[\text{Var}(x')/\text{Var}(x)]^{1/2}$	Mean (centroid) frequency of the signal; rises as energy shifts to higher frequencies within the band.
Hjorth complexity	$M(x')/M(x)$	Spectral spread; increases as the signal departs from a single dominant frequency toward broadband content.

Table 2: Time-domain features. Notation: N signal length; $\bar{x}$ sample mean; σ population standard deviation; $x_{\text{peak}} = (\max x - \min x)/2$ half the peak-to-peak range.

All frequency-domain features are computed from the one-sided FFT magnitude spectrum of the sweep. The spectral skewness and kurtosis are standardised by $\sigma_w^p \Sigma S$ and are therefore dimensionless and invariant to overall signal amplitude, so that spectral-shape information is separated from the amplitude information carried by RMS.

4.5 Candidate feature sets

Extracting the 14-feature set independently for the broadband signal and each sub-band yields $14 \times 3 = 42$ candidate AE features (broadband plus two sub-bands) and $14 \times$

Name		Formula	Description
Dominant frequency	fre-	$f_{\arg \max_k S_k^2}$	Location of the strongest spectral peak; tracks the most energetic resonance or excitation in the band.
Centre frequency	fre-	$f_c = \sum_k f_k S_k / \Sigma S$	Energy-weighted mean frequency; shifts as the balance between low- and high-frequency content changes.
Spectral width	band-	$\sigma_w = (\sum_k (f_k - f_c)^2 S_k / \Sigma S)^{1/2}$	Concentration of energy about the centroid; small for tonal signals, large for broadband activity.
Spectral skewness	skew-	$\sum_k (f_k - f_c)^3 S_k / (\sigma_w^3 \Sigma S)$	Direction of spectral asymmetry; indicates whether energy leans below or above the centroid.
Spectral kurtosis		$\sum_k (f_k - f_c)^4 S_k / (\sigma_w^4 \Sigma S)$	Tailedness of the spectrum; large when energy is spread out.
Spectral flatness		$\exp(K'^{-1} \sum_{k:S_k > 0} \ln S_k) / \bar{S}$	Tonal-versus-noise character; near 1 for broadband noise, near 0 for a pure tone.

Table 3: Frequency-domain features (companion to Table 2). Notation: $S_k = |X_k|$ one-sided FFT magnitude at bin k ; f_k corresponding frequency; $\bar{S}$ arithmetic mean of $\{S_k\}$; $\Sigma S = \sum_k S_k$; $f_c = (\Sigma S)^{-1} \sum_k f_k S_k$; $\sigma_w = ((\Sigma S)^{-1} \sum_k (f_k - f_c)^2 S_k)^{1/2}$.

4 = 56 candidate US features (broadband plus three sub-bands) per sweep. Features are named by band and statistic (e.g. `AE_500-1000kHz__mobility`); broadband features carry no band prefix. These two candidate sets, and their union (the combined AE+US set, 42 + 56 features), define the three sensor configurations—AE-only, US-only, and combined—evaluated in Sections 5–6. The reduction of these candidate sets to compact, non-redundant subsets for modelling is described in Section 5.

5 Feature Analysis

This section reduces the 42 AE and 56 US candidate features (Section 4) to compact, non-redundant subsets for modelling, and characterises what the surviving features respond to. The analysis proceeds in two selection stages – a relevance filter against κ followed by a redundancy reduction – after which the operating-condition dependence and low-dimensional structure of the retained sets are examined.

5.1 Two-stage feature selection

Stage 1 – relevance. Each candidate feature is ranked by its absolute Spearman ($|\rho|$, monotonic association) and Pearson ($|r|$, linear association) correlation with κ across all sweeps. Features must satisfy both $|\rho| \geq 0.1$ and $|r| \geq 0.1$ to be retained, removing features carrying essentially no monotonic or linear information about the target.

Stage 2 – redundancy. Among the survivors, multicollinearity is reduced by iterative variance-inflation-factor (VIF) elimination. At each round the VIF of every remaining feature is computed; if any feature exceeds the threshold $\text{VIF} = 5$, the *least κ -relevant* feature among those above the threshold (lowest $|\rho|$ with κ) is dropped, and all VIFs are recomputed on the reduced set. The procedure repeats until no remaining feature exceeds the threshold. Recomputing the VIFs each round addresses both pairwise and distributed (many-to-one) collinearity, while the relevance-based tie-break protects the strongest predictor in each collinear group: a highly inter-correlated feature is retained when it is the most κ -relevant member of its group, its less-relevant partners being removed in its place. By construction every retained feature has $\text{VIF} < 5$ on the final subset.

Table 4 summarises the selection funnel.

Table 4: Feature selection funnel: candidate features, survivors of the Stage 1 relevance filter, and the final retained sets after Stage 2 redundancy reduction.

Channel	Candidates	After Stage 1	Retained
AE	42	36	11
US	56	35	12
Combined	–	–	23

5.2 Correlation with κ

Table 5 lists the top ten features for each sensor ranked by combined Spearman–Pearson rank score with κ ; the complete rankings of all candidates, including Stage 2 retention, are given in Appendix A (Tables 9 and 10).

Table 5: Top 10 features by combined Spearman–Pearson rank score with κ , for the AE and the US signals. Features retained after Stage 2 redundancy reduction are marked with † .

AE			US		
Feature	$ \rho $	$ r $	Feature	$ \rho $	$ r $
500 kHz–1 MHz mobility †	0.848	0.662	10–20 kHz RMS	0.691	0.609
500 kHz–1 MHz centre freq.	0.799	0.687	RMS	0.614	0.547
Shape factor †	0.703	0.625	0–10 kHz RMS	0.610	0.542
500 kHz–1 MHz dominant freq.	0.676	0.582	10–20 kHz mobility †	0.583	0.536
500 kHz–1 MHz complexity	0.682	0.529	20–100 kHz mobility	0.420	0.466
20–500 kHz shape factor	0.631	0.554	20–100 kHz RMS	0.507	0.344
500 kHz–1 MHz RMS †	0.617	0.528	Complexity †	0.542	0.297
20–500 kHz centre freq.	0.563	0.543	20–100 kHz complexity †	0.416	0.396
Centre freq.	0.577	0.524	0–10 kHz kurtosis †	0.465	0.331
20–500 kHz mobility	0.532	0.550	Kurtosis	0.452	0.300

The top-ranked AE feature is 500 kHz–1 MHz mobility ($|\rho| = 0.848$), and the top-

ranked US feature is 10–20 kHz RMS ($|\rho| = 0.691$).

Two patterns distinguish the channels. For AE, relevance concentrates in the 500 kHz–1 MHz sub-band, which supplies five of the ten highest-ranked features (including the top-ranked mobility) and is led by spectral-shape and frequency-distribution descriptors—mobility, centre and dominant frequency, complexity—rather than by raw amplitude: where the high-frequency energy sits and how peaked its spectrum is carry more monotonic information about κ than its magnitude alone. For US, relevance instead concentrates in the low-frequency sub-bands (≤ 20 kHz) and is dominated by amplitude (RMS) features, with the broadband and 0–20 kHz RMS terms heading the ranking. The channels also differ in strength: the best AE feature exceeds the best US feature by a wide margin in $|\rho|$ (0.848 vs. 0.691), consistent with the AE sensor’s more direct mechanical coupling to asperity-contact events. Cross-referencing the retained features against operating conditions (Section 5.4) shows that the most κ -relevant AE features are also those most strongly correlated with rotational speed, with one informative exception: the 20–500 kHz skewness feature is almost uncorrelated with speed yet carries the strongest temperature association of the retained set, anticipating the within-step temperature channel examined below.

Broadband AE features rank lower than their sub-band counterparts, indicating that the frequency-specific content of the high-frequency AE signal carries more monotonic information about κ than the full-band aggregate [7, 22].

5.3 Redundancy structure

Figure 7 shows the variance inflation factors of the Stage 1 survivors, computed on the full survivor set, on a logarithmic axis. The values span several orders of magnitude. A feature plotted far above the threshold may still be retained: the iterative procedure (Section 5.1) removes its less κ -relevant collinear partners first, so its *recomputed* VIF on the reduced set falls below the threshold even though its full-set value is large.

The redundancy concentrates in the amplitude-ratio family. RMS, crest factor, margin factor, and shape factor are algebraically interrelated normalisations of the same waveform amplitude, and across sub-bands they form a tightly collinear cluster whose members carry the largest variance inflation factors—several orders of magnitude above the threshold. The iterative elimination therefore prunes this cluster most heavily, removing the duplicated band-copies and lower-relevance amplitude ratios until the surviving amplitude features are mutually non-collinear. A secondary, weaker cluster links the spectral-location descriptors (centre frequency, dominant frequency, bandwidth) within a band. By contrast, the distributional-shape features—skewness and spectral kurtosis—are close to mutually orthogonal and enter the retained set almost untouched. Because the candidate set is canonical (Section 4), these collinearities reflect genuine shared structure—primarily

the common amplitude scaling imposed by the operating point—rather than definitional identities among overlapping feature definitions.




Figure 7: Variance inflation factors of the Stage 1 survivors, computed on the full survivor set (logarithmic axis). Retained features in blue, removed features in grey; the dashed line marks the $VIF = 5$ threshold of the Stage 2 iterative elimination. Values are capped at 10^6 for display.

5.4 Operating-condition dependence

The raw correlation between signal features and κ is partially driven by the shared dependence on rotational speed and temperature: RPM has a Spearman correlation of $\rho = +0.75$ with κ and temperature has $\rho = -0.61$. Higher RPM reduces the required viscosity ν_1 (Eq. (2)) and therefore increases κ , while simultaneously increasing signal energy through more frequent contact events per second. Higher temperature reduces the actual viscosity ν and therefore reduces κ , while also affecting material damping and structural dynamics.

As a consequence, amplitude features such as RMS are partially acting as indirect proxies for rotational speed rather than as direct indicators of lubrication film adequacy [7, 17, 21]. This confounding is inherent to any single-operating-point sensor system without an explicit operating-condition normalisation step. Because κ is fully determined by the operating point (Section 2.1), a model given measured RPM and temperature directly would predict κ near-perfectly by construction; the contribution of the acoustic features is to recover this information without tachometer or thermometer access. This proxy structure is quantified directly. A LightGBM regressor trained on the retained AE features predicts RPM with $R^2 = 0.945$ (RMSE 233 rpm) and temperature with $R^2 = 0.663$ (RMSE 9.6 °C) on the hold-out set. Feeding these predicted operating points through Eqs. (1)–(2) yields a two-stage κ estimate with $R^2 = 0.826$ (RMSE 0.112), nearly matching the direct features-to- κ model ($R^2 = 0.834$, RMSE 0.109). The direct model is therefore

well described as an acoustic operating-point soft sensor combined with the ISO 281 mapping, with only a marginal contribution beyond it.

The features are not, however, mere speed trackers. Within a fixed staircase RPM step the drive state is approximately constant while the lubricant viscosity varies across the 13 temperature blocks; in this setting the strongest feature (20–500 kHz mobility) retains a median within-step Spearman correlation of -0.56 with temperature (predominantly negative across steps, i.e. tracking κ ; Figure 8), while remaining constant during the stationary heater-driven warm-up. The features therefore sense both κ channels—speed and temperature—through the rotating bearing, rather than reading either operating quantity directly.





Figure 8: Within-RPM-step Spearman correlation between selected AE sub-band features and temperature. At fixed speed the drive state is constant, so non-zero correlations reflect sensitivity to the lubrication state rather than to the operating point.

5.5 Low-dimensional structure and regime separability

Figure 9 shows the first two principal components of the retained AE and US feature sets coloured by κ value. In the AE projection the κ value varies smoothly across the plane rather than forming disjoint clusters: high- κ sweeps (well-separated films) concentrate in the upper-left lobe while low- κ sweeps fan out along a tail extending to large positive PC1, so the first two components—together capturing roughly two-thirds of the AE feature variance—already arrange the data along a continuous lubrication gradient. The US projection shows the same ordering more weakly: a single dense core carries a shallower κ gradient and a detached low- κ cluster sits apart from it, consistent with the lower predictive power of the US channel (Section 6). The continuity of the AE gradient, rather than sharp regime boundaries, mirrors the continuous nature of κ and indicates that the retained features encode a graded film-thickness signal rather than a small number of

Source the exact fraction of steps with negative within-step ρ from 09_proxy and add it as a macro; the previous “91%” referred to the now-removed 1–2 MHz complexity feature.

discrete operating states.



Confirm
pca_kappa
is ren-
dered
from
the final
visit-
grouped
run.

Figure 9: PCA projection of the retained AE and US feature sets, coloured by lubrication κ value.

6 Regression Modelling and Statistical Evaluation

6.1 Experimental design

6.1.1 Data sets and splits

The data experiments utilise three different features sets: the AE, the US, and the combined AE and US feature sets. Each of the later described models are fitted to each of the three datasets. Each dataset is divided into a training set (80%) and a hold-out test set (20%) by a *visit-grouped* split: sweeps are partitioned into groups corresponding to contiguous 60-second RPM holds (both sensor rows of a sweep kept together), and whole groups—rather than individual sweeps—are assigned to the training or hold-out set with `GroupShuffleSplit`. This prevents the near-duplicate windows of a single slowly-varying hold from straddling the train/test boundary, as a sweep-level random split would allow. The hold-out set is withheld entirely from all model development steps and used only for final evaluation. The training set is further divided during nested cross-validation described in Section 6.2.

6.1.2 Model families

Three regression model families are evaluated for three sensor configurations, as summarised in Table 6.

6.2 Repeated nested cross-validation

To obtain stable performance estimates that properly account for the variance introduced by the train/validation split, the model evaluation uses a repeated nested k -fold cross-

Model	Rationale	Hyperparameters & tuning
Elastic Net	$\ell_1 + \ell_2$ regularised linear model; ℓ_1 promotes sparsity, ℓ_2 stabilises correlated features.	α (9 values, 10^{-4} – 10^2) $\times$ ℓ_1 ratio (6 values, 0.1–1.0); ElasticNetCV , inner 5-fold.
Polynomial (deg. 2)	Augments linear model with all pairwise interaction terms; remains linear in parameters.	Ridge penalty λ (15 log-spaced values, 10^{-3} – 10^4); RidgeCV (leave-one-out within each outer fold).
LightGBM	Gradient-boosted trees; captures non-linear feature– κ mappings and high-order interactions.	Learning rate, num. leaves, max. depth, min. child samples, ℓ_2 reg.; Bayesian optimisation (Optuna TPE, 20 trials per outer fold).

Table 6: Regression model families: rationale and hyperparameter tuning strategy.

validation scheme. The outer loop is repeated $R = 5$ times with independently re-seeded, group-respecting 5-fold splits (the same hold-group structure as the hold-out split, so near-duplicate windows never straddle a fold boundary), yielding $R \times k = 25$ outer RMSE scores per model–sensor combination. See Figure 10 for data partitions.




Figure 10: Data partitioning for the repeated nested cross-validation: an 80/20 group-level hold-out split (groups = 60-second RPM holds), with the training pool further divided into $R = 5$ independently seeded, group-respecting 5-fold outer splits and inner hyperparameter-selection folds.

For LightGBM, each outer fold contains its own inner Optuna optimisation loop,

ensuring that the reported outer scores are not contaminated by hyperparameter selection on the outer fold data. For the linear models, hyperparameter selection runs within each outer training partition: `ElasticNetCV` performs an inner 5-fold search over the α/ℓ_1 -ratio grid, and `RidgeCV` a leave-one-out generalised cross-validation over the λ grid; the selected hyperparameters are then refit on the full outer-training fold before the outer validation fold is scored.

The 25 outer scores form the empirical sampling distribution of RMSE for each model and are used directly in the significance tests described below.

6.3 Statistical significance testing

Pairwise significance tests are conducted at two levels.

Level 1 — Within-sensor comparisons (architecture differences). For pairs of models trained on the same sensor configuration (e.g., Elastic Net vs. Polynomial on AE), the corrected repeated k -fold t -test of Nadeau and Bengio [29] is applied to the 25 paired RMSE scores. This test corrects the standard paired t -test for the positive correlation induced by shared training data across repeated folds. Multiple within-sensor comparisons are adjusted by the Holm–Bonferroni step-down procedure at $\alpha = 0.05$.

Level 2 — Cross-sensor comparisons (best per sensor configuration). For the best model from each sensor configuration (selected by lowest mean CV RMSE), pairwise differences on the common hold-out set are assessed by:

- The Wilcoxon signed-rank test on paired residuals (non-parametric);
- The modified Diebold–Mariano test for equal predictive accuracy under squared-error loss;
- A bootstrap Δ RMSE 95% confidence interval (10 000 resamples with replacement).

Cross-sensor p -values are also corrected by Holm–Bonferroni.

6.4 Results

Model	Sensors	CV RMSE	HO R^2	HO MAE	HO RMSE
LightGBM	Combined	0.0876 ± 0.0203	0.908	0.056	0.081
LightGBM	AE	0.1138 ± 0.0177	0.834	0.082	0.109
LightGBM	US	0.1648 ± 0.0372	0.653	0.105	0.158
Polynomial	Combined	0.1848 ± 0.0416	0.677	0.116	0.153
Elastic Net	Combined	0.1913 ± 0.0434	0.593	0.134	0.171
Polynomial	AE	0.2027 ± 0.0381	0.487	0.155	0.192
Polynomial	US	0.2236 ± 0.0552	0.333	0.163	0.219
Elastic Net	AE	0.2268 ± 0.0455	0.491	0.152	0.192
Elastic Net	US	0.2279 ± 0.0525	0.463	0.149	0.197

Table 7: Regression model performance. CV RMSE: mean $\pm$ std of outer RMSE across 25 repeated nested CV folds (training set $n = 7\,178$). HO: hold-out test set ($n = 1\,326$). Best hold-out R^2 in bold.

Overall performance. LightGBM achieves the best hold-out performance across all sensor configurations, with AE-only reaching $R^2 = 0.834$, RMSE = 0.109 κ -units. On a target range of $[0.04, 1.64]$ with mean 0.51, this corresponds to a relative RMSE of approximately 21% of the mean target value. Linear and polynomial models on AE features achieve $R^2 = 0.487$ –0.491, confirming a non-linear feature- κ mapping that gradient boosting exploits more effectively.

Sensor comparison. AE features substantially outperform US features across all model families. The LightGBM AE-US gap ($\Delta R^2 = 0.18$) is large and significant on the hold-out set (Wilcoxon $p = 1.7 \times 10^{-7}$; Δ RMSE = 0.0487, 95% CI $[0.0368, 0.0612]$), although the corrected repeated nested-CV t -test does not reach significance (Nadeau-Bengio $p = 0.121$); we report both. This reflects the richer frequency content and more direct physical sensitivity of the AE sensor to asperity contact events.

Sensor fusion. LightGBM on the combined AE+US feature set achieves $R^2 = 0.908$ ($\Delta R^2 = 0.074$ over the AE-only result of 0.834). The within-CV corrected t -test gives $p = 5.1 \times 10^{-6}$; on the hold-out set, Wilcoxon $p = 9.6 \times 10^{-39}$, Diebold-Mariano $p = 6.5 \times 10^{-18}$, and the bootstrap Δ RMSE = 0.0281 with 95% CI $[0.0219, 0.0343]$. For the linear models, the combined feature set performs similarly to AE alone (Elastic Net: R^2 0.593 vs. 0.491; Polynomial: 0.677 vs. 0.487). The fusion gain is small in absolute terms but consistent across all four tests, indicating that the US channel contributes a modest amount of non-redundant information.

Model consistency. The close agreement between nested CV RMSE and hold-out RMSE across all models indicates stable generalisation with no evidence of significant overfitting.

Figure 11 shows predicted vs. actual κ for the hold-out set for all model families under the AE-only configuration.




Figure 11: Predicted vs. actual κ on the hold-out test set for all model families and sensor configurations, colored according to RPM. The dashed diagonal indicates perfect prediction. Residual spread increases in the mixed lubrication regime ($0.5 \leq \kappa < 1.0$).

SHAP-based feature importance. SHAP (SHapley Additive exPlanations) values computed via TreeExplainer identify which AE sub-band features drive the κ predictions. For LightGBM the dominant contributor by a large margin is *500 kHz–1 MHz mobility* (mean $|\text{SHAP}| = 0.205$), followed by *20–500 kHz RMS* (0.050) and *500 kHz–1 MHz crest factor* (0.036). The 500 kHz–1 MHz mobility feature ranks *first* for both LightGBM and the polynomial model and third for Elastic Net (behind the broadband shape factor and 20–500 kHz RMS), making it the most robust single predictor of κ across model families (cross-model agreement, Table 8). Figure 12 shows the full importance ranking.

Update
visual
interpretation
of the figure
after new
results.

Feature	ElasticNet	Polynomial	LightGBM	Models in top 10
500 kHz–1 MHz mobility	3	1	1	3
Broadband shape factor	1	4	4	3
Broadband margin factor	5	2	6	3
500 kHz–1 MHz crest factor	4	10	3	3
500 kHz–1 MHz RMS	7	9	5	3
20–500 kHz RMS	2	–	2	2
20–500 kHz skewness	6	–	7	2
500 kHz–1 MHz skewness	9	–	8	2
500 kHz–1 MHz spectral kurtosis	8	–	9	2

Table 8: Cross-model SHAP agreement for AE features: features appearing in the top 10 by mean $|\text{SHAP}|$ for more than one model type. Polynomial ranks are positions in the full expanded-term importance list (interaction and squared terms included); only base features are tabulated.




Figure 12: SHAP feature importance (mean $|\text{SHAP value}|$) for the LightGBM model trained on AE features. Features from frequency sub-bands are prefixed with the band label (e.g., AE_20-500kHz_).

7 Discussion

7.1 AE vs. US

The consistently superior performance of AE over the US channel across all model families reflects a fundamental hardware difference. The AE sensor captures wideband, rotation-generated content extending into the MHz range (analysed here up to 1 MHz), dominated by verified stress-wave emission up to several hundred kHz. The US channel down-mixes its $\approx 38\text{--}40$ kHz sensing band to an audio-frequency baseband (Section 2.2). Although

the probe acquires content across its 0–100 kHz band (Section 4.2), the lubrication-relevant information is concentrated in the nominal 0–20 kHz baseband—a mixture of down-mixed contact activity, low-frequency machinery noise, VFD harmonics, and structural vibration, within which the 10–20 kHz half carries the strongest κ correlations (Section 5.2)—while the 20–100 kHz above-baseband content contributes comparatively little. The high-frequency transient signature of individual microscale contact events, which the AE channel captures directly, is largely absent from this down-mixed envelope. Replacing the heterodyned probe with a wideband passive acoustic or ultrasound sensor would be expected to substantially improve the US-channel models.

7.2 Fusion gain

The fusion gain ($\Delta R^2 = 0.074$, $\Delta \text{RMSE} = 0.0281$ on the hold-out) is small but statistically robust, detected under repeated nested cross-validation and confirmed by all three hold-out tests. Its limited magnitude can be understood from two complementary angles. First, the US channel retains only 12 features after selection (vs. 11 for AE), and the top US correlations ($|\rho| \leq 0.691$) are weaker than the top AE correlations ($|\rho| = 0.848$). The US channel therefore contributes a genuine but limited amount of non-redundant information to the combined model. Second, the US channel’s down-mixed audio-frequency baseband means its information content partially overlaps with the low-frequency tail of the AE signal, limiting the independent contribution it can make.

This result is an informative contribution to the LCM literature, where AE–ultrasound complementarity is widely *argued* but rarely *tested*. Whether a wideband passive ultrasound sensor (rather than a heterodyned probe) would provide more independent information – and thus a more substantial fusion gain – remains an open question for future work.

7.3 Interpreting the amplitude-confounding result

The strong raw Spearman correlations of amplitude features (RMS and related statistics) and frequency-tracking features (Hjorth mobility, centre frequency) with κ should be interpreted with caution. As discussed in Section 5.4, rotational speed is a driver of both κ (through ν_1 in Eq. (2)) and broadband signal energy (through the rate of and energy in contact events per revolution as described in [7]). A model trained on amplitude features therefore exploits speed as an indirect proxy for κ rather than sensing purely lubrication-specific signal changes directly.

This confounding does not invalidate the regression results: a model that uses speed-mediated amplitude changes to predict κ is still predicting κ , and its predictions may be operationally useful. However, it does indicate that the model’s generalisation to operating conditions outside its training distribution (different speed ranges, different

bearing loads) may be limited. The two-stage analysis of Section 5.4 makes this concrete: predicted operating points alone, fed through the ISO 281 mapping, recover $R^2 = 0.826$ of the direct model’s $R^2 = 0.834$ —the model is, to first order, an operating-point soft sensor, and its lubrication sensitivity beyond the operating point rests on the within-step evidence of Section 5.4. Future feature engineering should prioritise amplitude-invariant spectral-shape features (spectral skewness and kurtosis, spectral bandwidth, spectral flatness) that are expected to carry a higher proportion of genuine lubrication-state information relative to operating condition proxies.

7.4 Implications for the κ regression framework

The κ regression framework is validated as a methodologically sound approach for passive non-invasive LCM, when κ is assumed to be representative of the lubrication conditions in the bearing. A hold-out $R^2 = 0.834$ and $\text{RMSE} = 0.109$ κ -units for a single AE sensor represents strong performance across a demanding continuous speed-and-temperature ramp spanning boundary through full-film lubrication conditions. Under matched random (sweep-level) splits this performance is on par with the passive κ regression reported by Jakobsen et al. [1, 16]. The figures reported here are not directly comparable to most prior work, however, because they are obtained under a stricter *visit-grouped* evaluation protocol (Section 6.1) that removes the near-duplicate leakage admitted by a sweep-level random split; to our knowledge they are the first κ -regression results reported under such a leakage-controlled protocol. The contribution is thus methodological—a more conservative, deployment-representative performance estimate—rather than a raw accuracy gain over earlier studies.

The repeated nested cross-validation design provides appropriately conservative performance estimates: the 25 outer RMSE scores capture fold-to-fold variance without the optimistic bias introduced by simple k -fold evaluation with shared hyperparameter search. The corrected significance tests confirm that performance differences between model families are assessed under assumptions that properly account for data re-use.

7.5 Limitations

1. **Single bearing and lubricant.** All results are specific to the SYJ25 bearing and Keratech 22 oil. Generalisability to other bearing types, sizes, or lubricants is unknown.
2. **Heterodyned US channel.** The passive US probe down-mixes a narrow ≈ 38 – 40 kHz sensing band to an audio-frequency baseband rather than recording wideband ultrasound directly, and its lubrication-relevant information is concentrated below ≈ 20 kHz (Section 2.2). It should not be equated with a wideband passive

ultrasound sensor, and the US-channel results should be interpreted in light of this hardware constraint.

3. **Incomplete operating-condition coverage.** The sampled operating points do not form a full speed–temperature factorial: high rotational speeds are reached predominantly at lower temperatures and high temperatures predominantly at lower speeds, leaving the high-speed/high-temperature corner sparsely populated (Figure 2). The κ distribution is correspondingly skewed toward the boundary and mixed-lubrication regime, with full-film states ($\kappa > 1$) under-represented. Extrapolation into the under-sampled region of the operating space should therefore be treated with caution.
4. **Model-derived κ reference.** The viscosity ratio is not measured directly but computed for each sweep from the measured operating point—rotational speed and housing temperature—via the ISO 281 relations and the ASTM D341 Walther viscosity–temperature equation (Section 2.1). The regression target is therefore a model-derived reference that inherits the assumptions of those relations, notably the use of housing temperature as a proxy for the elastohydrodynamic contact temperature; this also underlies the operating-point soft-sensor interpretation of Section 7.3.

8 Conclusion

This paper has presented a controlled study of simultaneous acoustic emission and passive ultrasound sensing regressed to the ISO 281 viscosity ratio κ for non-invasive lubrication condition monitoring of a ball bearing. The main findings are as follows:

1. **AE sub-band features are strong predictors of κ .** LightGBM on 11 retained AE features achieves hold-out $R^2 = 0.834$, RMSE = 0.109 κ -units. The 500 kHz–1 MHz mobility feature is the dominant SHAP contributor and the most robust single predictor across model families—ranked first for the LightGBM and polynomial models (third for Elastic Net)—identifying the upper AE sub-band (500 kHz–1 MHz) as the most informative frequency range for lubrication-regime prediction.
2. **Sensor fusion yields a small but consistently significant gain.** The LightGBM combined model achieves $R^2 = 0.908$ ($\Delta R^2 = 0.074$ over AE alone); within-CV corrected t -test $p = 5.1 \times 10^{-6}$; hold-out tests: Wilcoxon $p = 9.6 \times 10^{-39}$, Diebold–Mariano $p = 6.5 \times 10^{-18}$, bootstrap 95% CI [0.0219, 0.0343] excluding zero. The heterodyned US channel thus supplies a modest amount of non-redundant information, though its 12 retained features are individually weaker than the AE features.

3. **AE substantially outperforms the US channel.** The AE-US gap ($\Delta R^2 = 0.18$ for LightGBM) is large and significant on the hold-out tests (Wilcoxon $p = 1.7 \times 10^{-7}$), although the corrected repeated-CV t -test does not reach significance ($p = 0.121$); we therefore report both. The gap reflects the bandwidth limitation of the heterodyned US probe.
4. **Rigorous evaluation via repeated nested cross-validation.** The 25 outer RMSE scores per model provide stable, conservative performance estimates; corrected significance tests confirm that all pairwise model differences within each sensor configuration are significant after Holm-Bonferroni correction, with LightGBM strongest throughout.
5. **The κ regression framework is workable for passive non-invasive LCM, with a quantified caveat.** Most of the headline performance is recoverable from acoustically inferred operating points alone ($R^2 = 0.826$ via the two-stage decomposition), so the framework is best understood as an acoustic operating-point soft sensor combined with the ISO 281 mapping; the within-speed-step analysis nevertheless confirms a genuine, temperature-mediated film channel beyond the operating-point proxy.

Future work should address the identified limitations by: validating the model-derived κ reference against an independent film-thickness proxy, using slipring-based conductance measurements on the same test rig (data collection ongoing); acquiring wideband passive ultrasound to replace the heterodyned probe; validating across multiple independent sessions, bearing types, and lubricants; investigating temporal autocorrelation and its effect on hold-out validity; and developing physically motivated sub-band ratio features that are explicitly designed to be robust to operating-condition variation.

A Full Feature Correlation Rankings

Tables 9 and 10 list the complete Spearman and Pearson correlations with κ for all AE and US features that passed the initial data quality filter, sorted by combined rank. Features retained after the Stage 2 redundancy reduction are indicated with a checkmark.

Table 9: Full AE feature correlation ranking with κ . $|\rho|$: absolute Spearman correlation; $|r|$: absolute Pearson correlation. Retained: selected after Stage 2 redundancy reduction.

Rank	Sub-band	Feature	$ \rho $	$ r $	Ret.
1	500–1000kHz	mobility	0.848	0.662	✓
2	500–1000kHz	center frequency	0.799	0.687	–

(continued)

Rank	Sub-band	Feature	$ \rho $	$ r $	Ret.
3	Broadband	shape factor	0.703	0.625	✓
4	500–1000kHz	dominant frequency	0.676	0.582	–
5	500–1000kHz	complexity	0.682	0.529	–
6	20—500 kHz	shape factor	0.631	0.554	–
7	500–1000kHz	rms	0.617	0.528	✓
8	20—500 kHz	center frequency	0.563	0.543	–
9	Broadband	center frequency	0.577	0.524	–
10	20—500 kHz	mobility	0.532	0.550	–
11	500–1000kHz	crest factor	0.553	0.523	✓
12	Broadband	margin factor	0.589	0.485	✓
13	Broadband	mobility	0.533	0.535	–
14	500–1000kHz	margin factor	0.509	0.512	–
15	Broadband	complexity	0.566	0.454	–
16	20—500 kHz	complexity	0.523	0.473	–
17	20—500 kHz	rms	0.603	0.339	✓
18	Broadband	rms	0.601	0.320	–
19	20—500 kHz	dominant frequency	0.381	0.521	–
20	500–1000kHz	kurtosis	0.469	0.404	–
21	20—500 kHz	margin factor	0.464	0.343	–
22	Broadband	crest factor	0.504	0.288	–
23	Broadband	kurtosis	0.491	0.174	–
24	20—500 kHz	spectral kurtosis	0.390	0.267	–
25	20—500 kHz	spectral skewness	0.372	0.220	✓
26	500–1000kHz	shape factor	0.296	0.289	–
27	500–1000kHz	skewness	0.364	0.198	✓
28	500–1000kHz	spectral kurtosis	0.298	0.258	✓
29	20—500 kHz	crest factor	0.371	0.157	–
30	20—500 kHz	skewness	0.209	0.261	✓
31	Broadband	spectral skewness	0.284	0.174	✓
32	20—500 kHz	spectral flatness	0.246	0.165	–
33	20—500 kHz	spectral bandwidth	0.241	0.164	–
34	500–1000kHz	spectral skewness	0.198	0.157	–
35	Broadband	spectral kurtosis	0.229	0.108	–
36	500–1000kHz	spectral bandwidth	0.135	0.105	–

Table 10: Full US feature correlation ranking with κ . $|\rho|$: absolute Spearman correlation; $|r|$: absolute Pearson correlation. Retained: selected after Stage 2 redundancy reduction.

Rank	Sub-band	Feature	$ \rho $	$ r $	Ret.
1	10–20kHz	rms	0.691	0.609	–
2	Broadband	rms	0.614	0.547	–
3	0–10kHz	rms	0.610	0.542	–
4	10–20kHz	mobility	0.583	0.536	✓
5	20–100kHz	mobility	0.420	0.466	–
6	20–100kHz	rms	0.507	0.344	–
7	Broadband	complexity	0.542	0.297	✓
8	20–100kHz	complexity	0.416	0.396	✓
9	0–10kHz	kurtosis	0.465	0.331	✓
10	Broadband	kurtosis	0.452	0.300	–
11	10–20kHz	crest factor	0.350	0.394	✓
12	0–10kHz	mobility	0.380	0.348	✓
13	10–20kHz	complexity	0.342	0.306	✓
14	10–20kHz	margin factor	0.280	0.358	–
15	20–100kHz	dominant frequency	0.277	0.360	–
16	Broadband	dominant frequency	0.322	0.306	–
17	0–10kHz	dominant frequency	0.322	0.304	✓
18	Broadband	margin factor	0.298	0.309	–
19	Broadband	crest factor	0.309	0.290	–
20	10–20kHz	kurtosis	0.232	0.347	–
21	0–10kHz	margin factor	0.260	0.289	–
22	0–10kHz	crest factor	0.269	0.253	–
23	Broadband	shape factor	0.238	0.280	–
24	20–100kHz	spectral bandwidth	0.267	0.245	✓
25	Broadband	skewness	0.298	0.206	✓
26	20–100kHz	spectral flatness	0.255	0.234	–
27	0–10kHz	shape factor	0.225	0.254	–
28	0–10kHz	skewness	0.276	0.190	–
29	10–20kHz	shape factor	0.113	0.305	–
30	20–100kHz	crest factor	0.184	0.230	–
31	20–100kHz	spectral skewness	0.263	0.144	✓
32	20–100kHz	center frequency	0.176	0.187	–
33	10–20kHz	spectral bandwidth	0.167	0.136	✓
34	10–20kHz	spectral flatness	0.164	0.112	–

CRediT authorship contribution statement

Frederik Appel Vardinghus-Nielsen: Conceptualization, Data curation, Formal analysis, Investigation, Methodology, Project administration, Resources, Software, Validation, Visualization, Writing – original draft, Writing – review & editing.

Declaration of competing interests

The author declares that there are no competing interests.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

IFD?

Data availability

The reproducible Python analysis pipeline and extracted features (Parquet format) accompany this paper.

Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work the author used AI-assisted writing tools for language editing and structural drafting. After using these tools, the author reviewed and edited the content as necessary and takes full responsibility for the content of the publication.

Acknowledgements

The author thanks Morten and colleagues at CeramicSpeed (Holstebro, Denmark) for providing the test facility, experimental assistance, and preliminary analysis.

Public release of the complete dataset (HDF5 signal files) is not yet decided—resolve with CeramicSpeed and choose: (a) full Zenodo deposit, (b) features + code

References

- [1] Morten Opprud Jakobsen, Eskild Sune Herskind, Christian Fischer Pedersen, and Mikael Bergholz Knudsen. Detecting insufficient lubrication in rolling bearings, using a low cost MEMS microphone to measure vibrations. *Mechanical Systems and Signal Processing*, 200:110553, 2023.
- [2] Wei Dai. In Situ Measurement of Grease Capacitive Film Thickness in Bearings: A Review. *Lubricants*, 12(10):329, 2024.
- [3] Mohamed G. A. Nassef, Mina Soliman, Belal Galal Nassef, Mohamed A. Daha, and Galal A. Nassef. Impact of Graphene Nano-Additives to Lithium Grease on the Dynamic and Tribological Behavior of Rolling Bearings. *Lubricants*, 10(2):29, 2022.
- [4] Dong Wang, Kwok-Leung Tsui, and Qiang Miao. Prognostics and health management: A review of vibration based bearing and gear health indicators. *IEEE Access*, 6:665–676, 2018.
- [5] Fei Shang, Bo Sun, and Dandan Cai. A comprehensive review of non-destructive methods for bearing lubricating oil film thickness measurement. *Industrial Lubrication and Tribology*, 76(3):441–451, 2024.
- [6] Matías Marticorena and Oscar García Peyrano. Rolling bearing condition monitoring technique based on cage rotation analysis and acoustic emission. *Journal of Dynamics, Monitoring and Diagnostics*, 1:57–65, 2022.
- [7] Xuanrui Zhang, Dong Zhen, Guojin Feng, Xiaoxia Liang, Hao Zhang, and Fengshou Gu. A coupled vibration-acoustic emission model for roller bearings incorporating lubricant traction and sliding friction. *Tribology International*, 221:112013, September 2026.
- [8] Shan Pang. Stacked maximum independence autoencoders: A domain generalization approach for fault diagnosis under various working conditions. *Mechanical Systems and Signal Processing*, 208:111035, 2024.
- [9] Zehui Hua, Juanjuan Shi, and Patrick Dumond. Domain-invariant feature exploration for intelligent fault diagnosis under unseen and time-varying working conditions. *Mechanical Systems and Signal Processing*, 224:112193, 2025.
- [10] Fei Jiang, Yicong Kuang, Tao Li, Shaohui Zhang, Zhaoqian Wu, Ke Feng, and Weihua Li. Towards enhanced interpretability: A mechanism-driven domain adaptation model for bearing fault diagnosis across operating conditions. *Mechanical Systems and Signal Processing*, 225:112244, 2025.

- [11] James M. Wakiru, Liliane Pintelon, Peter N. Muchiri, and Peter K. Chemweno. A review on lubricant condition monitoring information analysis for maintenance decision support. *Mechanical Systems and Signal Processing*, 118:108–132, 2019.
- [12] Daniel Cornel, Francisco Gutiérrez Guzmán, Georg Jacobs, and Stephan Neumann. Condition monitoring of roller bearings using acoustic emission. *Wind Energy Science*, 6(2):367–376, 2021.
- [13] Xinzhuo Zhang, Zhijun Ren, Kai Huang, Tanta Lin, Yongsheng Zhu, Hong Fu, and Jun Hong. An interpretable rolling-bearing lubrication-state diagnosis method based on multiscale residual and semantic consistency CycleGAN. *IEEE Transactions on Instrumentation and Measurement*, 75:3512116, 2026.
- [14] Jiayi Sun, Liming Wang, Jianfeng Li, Fangyi Li, Jianyong Li, and Haiyang Lu. Online oil debris monitoring of rotating machinery: A detailed review of more than three decades. *Mechanical Systems and Signal Processing*, 149:107341, 2021.
- [15] Jeremy Lineham. Ultrasonic probes for inspecting bearings. *World Pumps*, 2008(503):34–36, 2008.
- [16] Morten Opprud Jakobsen. Low cost MEMS accelerometer and microphone based condition monitoring sensor, with LoRa and Bluetooth Low Energy radio. *HardwareX*, 18:e00525, 2024.
- [17] J Miettinen, P Andersson, and V Wikström. Analysis of grease lubrication of a ball bearing using acoustic emission measurement. *Proceedings of the Institution of Mechanical Engineers, Part J: Journal of Engineering Tribology*, 215(6):535–544, 2001.
- [18] Takeo Yoshioka and Shigeo Shimizu. Monitoring of ball bearing operation under grease lubrication using a new compound diagnostic system detecting vibration and acoustic emission. *Tribology Transactions*, 52(6):725–730, 2009.
- [19] N. Tandon, K. M. Ramakrishna, and G. S. Yadava. Condition monitoring of electric motor ball bearings for the detection of grease contaminants. *Tribology International*, 40:29–36, 2007.
- [20] Cuiping Wang, Hongyuan Qi, Dongming Hou, Defu Han, and Jiangtian Yang. Fingerprint feature and dynamic threshold mechanism based on acoustic emission for bearing fault detection. *Mechanical Systems and Signal Processing*, 199:110496, 2023.
- [21] Geng Hou and Liangchi Zhang. In-situ detection of lubrication status of bearings with acoustic emission monitoring. *Wear*, 571:205839, 2025.

- [22] Philipp Renhart, Michael Maier, Christopher Strabegg, Florian Summer, Florian Grün, and Andreas Eder. Monitoring tribological events by acoustic emission measurements for bearing contacts. *Lubricants*, 9(11):109, 2021.
- [23] Jing Liu, Y. J. Xu, and Gang Pan. A combined acoustic and dynamic model of a defective ball bearing. *Journal of Sound and Vibration*, 501:116029, 2021.
- [24] Ling Wang and Honor Powrie. Advanced techniques for bearing condition monitoring. In *Electric Vehicle Tribology*, pages 153–181. Elsevier, 2024.
- [25] ISO 281:2007 — rolling bearings — dynamic load ratings and rating life, 2007.
- [26] Haiyang Pan, Yu Yang, Xin Li, Jinde Zheng, and Junsheng Cheng. Symplectic geometry mode decomposition and its application to rotating machinery compound fault diagnosis. *Mechanical Systems and Signal Processing*, 114:189–211, 2019.
- [27] Xinzhuo Zhang, Xuhua Zhang, Linbo Zhu, Chuang Gao, Bo Ning, and Yongsheng Zhu. A graph-data-based monitoring method of bearing lubrication using multi-sensor. *Lubricants*, 12(6):229, 2024.
- [28] Bo Hjorth. EEG analysis based on time domain properties. *Electroencephalography and Clinical Neurophysiology*, 29(3):306–310, 1970.
- [29] Claude Nadeau and Yoshua Bengio. Inference for the generalization error. *Machine Learning*, 52(3):239–281, 2003.